

**VALENTINE'S DAY SPECIAL: Up to 33% OFF** selected dedicated servers and components!

GET COUPON CODE



## Prerequisites

- An Ubuntu system.
- Access to a terminal or command line.
- A user with `sudo` or **root** permissions.

## Install Packages Required for Spark

Before downloading and setting up Spark, you need to install necessary dependencies. This step includes installing the following packages:

- JDK
- Scala
- Git

Open a terminal window and run the following command to install all three packages at once:

```
sudo apt install default-jdk scala git -y
```

You will see which packages will be installed.

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

```
Reading state information... Done
git is already the newest version (1:2.17.1-1ubuntu0.5).
The following packages were automatically installed and are no longer required:
  liballegro4.4 libdevilc2 libevent-core-2.1-6 libllvm7 liblua5.1-2
  liblua5.1-common libmng2 libmodplug1 libopenal-data libopenal1
  libphysfs1 libSDL1.2debian libSDL2-2.0-0 vim-runtime
Use 'sudo apt autoremove' to remove them.
The following additional packages will be installed: ←
  ca-certificates-java default-jdk-headless default-jre default-jre-headless
  fonts-dejavu-extra java-common libatk-wrapper-java libatk-wrapper-java-jni
  libhawtjni-runtime-java libice-dev libjansi-java libjansi-native-java
  libjline2-java libpthreads0-dev libsm-dev libx11-dev libx11-doc
  libxau-dev libxcb1-dev libxdmcp-dev libxt-dev openjdk-11-jdk
  openjdk-11-jdk-headless openjdk-11-jre openjdk-11-jre-headless
  scala-library scala-parser-combinators scala-xml x11proto-core-dev
  x11proto-dev xorg-sgml-doctools xtrans-dev
```

Once the process completes, **verify the installed dependencies** by running these commands:

```
java -version; javac -version; scala -version; git --version
```

```
test@ubuntu1:~$ java -version; javac -version; scala -version; git --version
openjdk version "11.0.6" 2020-01-14
OpenJDK Runtime Environment (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1)
OpenJDK 64-Bit Server VM (build 11.0.6+10-post-Ubuntu-1ubuntu118.04.1, mixed mode, sharing)
javac 11.0.6
Scala code runner version 2.11.12 -- Copyright 2002-2017, LAMP/EPFL
git version 2.17.1
```

The output prints the versions if the installation completed successfully for all packages.

## Download and Set Up Spark on Ubuntu

Now, **you need to download the version of Spark you want** from their website. We will go for *Spark 3.0.1* with *Hadoop 2.7* as it is the latest version at the time of writing this article.

Use the **wget** command and the direct link to download the Spark archive:

```
wget https://downloads.apache.org/spark/spark-3.0.1/spark-3.0.1-bin-hadoop2.7.tgz
```

When the download completes, you will see the *saved* message.

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

```
::2
Connecting to downloads.apache.org (downloads.apache.org)[88.99.95.219]:443... connected
.
HTTP request sent, awaiting response... 200 OK
Length: 219929956 (210M) [application/x-gzip]
Saving to: 'spark-3.0.1-bin-hadoop2.7.tgz.1'

spark-3.0.1-bin-hadoo 100%[=====>] 209.74M  2.78MB/s  in 53s
2020-09-14 19:22:16 (3.96 MB/s) - 'spark-3.0.1-bin-hadoop2.7.tgz.1' saved [219929956/219929956]
```



**Note:** If the URL does not work, please go to the [Apache Spark](#) download page to check for the latest version. Remember to replace the Spark version number in the subsequent commands if you change the download URL.

Now, extract the saved archive [using tar](#):

```
tar xvf spark-*
```

Let the process complete. The output shows the files that are being unpacked from the archive.

Finally, move the unpacked directory *spark-3.0.1-bin-hadoop2.7* to the ***opt/spark*** directory.

Use the **mv** command to do so:

```
sudo mv spark-3.0.1-bin-hadoop2.7 /opt/spark
```

The terminal returns no response if it successfully moves the directory. If you mistype the name, you will get a message similar to:

```
mv: cannot stat 'spark-3.0.1-bin-hadoop2.7': No such file or directory.
```

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

Use the **echo** command to add these three lines to *.profile*:

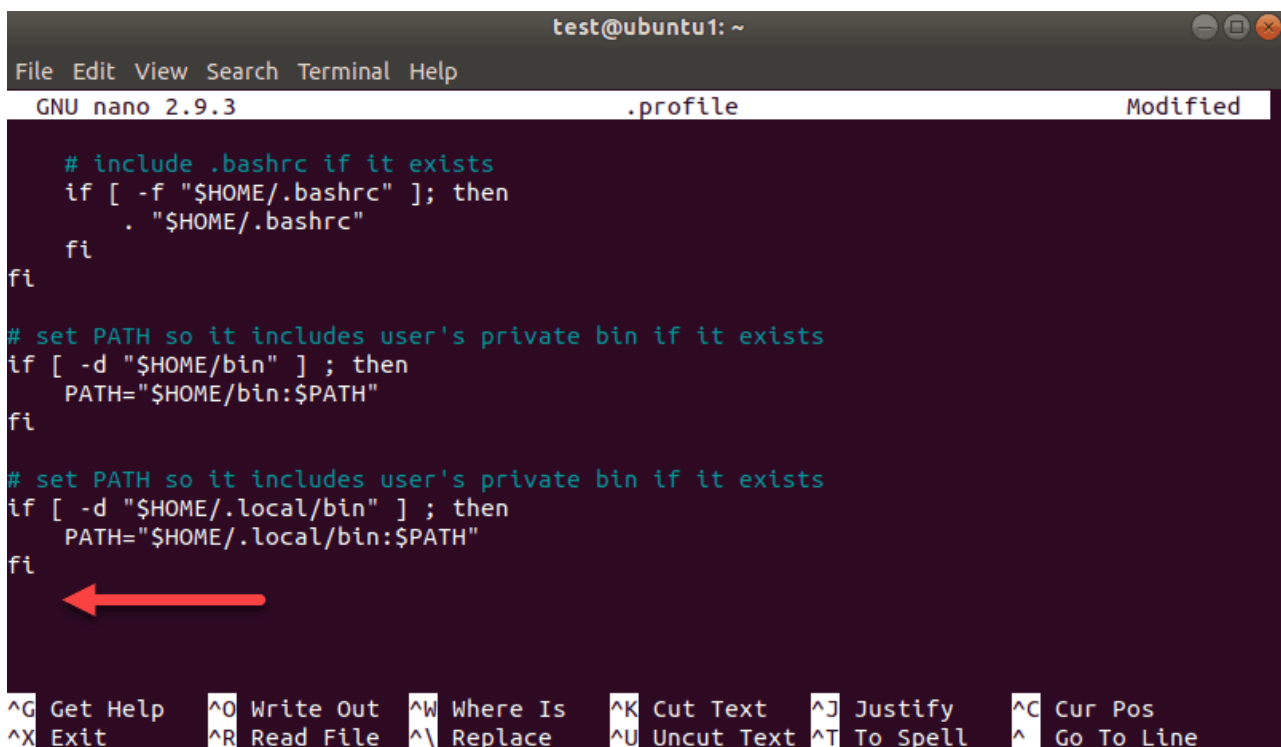
```
echo "export SPARK_HOME=/opt/spark" >> ~/.profile
echo "export PATH=$PATH:$SPARK_HOME/bin:$SPARK_HOME/sbin" >> ~/.profile
echo "export PYSPARK_PYTHON=/usr/bin/python3" >> ~/.profile
```

You can also add the export paths by editing the *.profile* file in the editor of your choice, such as nano or vim.

For example, to use nano, enter:

```
nano .profile
```

When the profile loads, scroll to the bottom of the file.



```
test@ubuntu1: ~
File Edit View Search Terminal Help
GNU nano 2.9.3 .profile Modified

# include .bashrc if it exists
if [ -f "$HOME/.bashrc" ]; then
    . "$HOME/.bashrc"
fi

# set PATH so it includes user's private bin if it exists
if [ -d "$HOME/bin" ] ; then
    PATH="$HOME/bin:$PATH"
fi

# set PATH so it includes user's private bin if it exists
if [ -d "$HOME/.local/bin" ] ; then
    PATH="$HOME/.local/bin:$PATH"
fi
```

←

^G Get Help ^O Write Out ^W Where Is ^K Cut Text ^J Justify ^C Cur Pos  
^X Exit ^R Read File ^\ Replace ^U Uncut Text ^T To Spell ^\_ Go To Line

Then, add these three lines:

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

```
export PYSPARK_PYTHON=/usr/bin/python3
```

Exit and save changes when prompted.

When you finish adding the paths, load the *.profile* file in the command line by typing:

```
source ~/.profile
```



## Start Standalone Spark Master Server

Now that you have completed configuring your environment for Spark, you can start a master server.

In the terminal, type:

```
start-master.sh
```



To view the Spark Web user interface, open a web browser and enter the localhost IP address on port 8080.

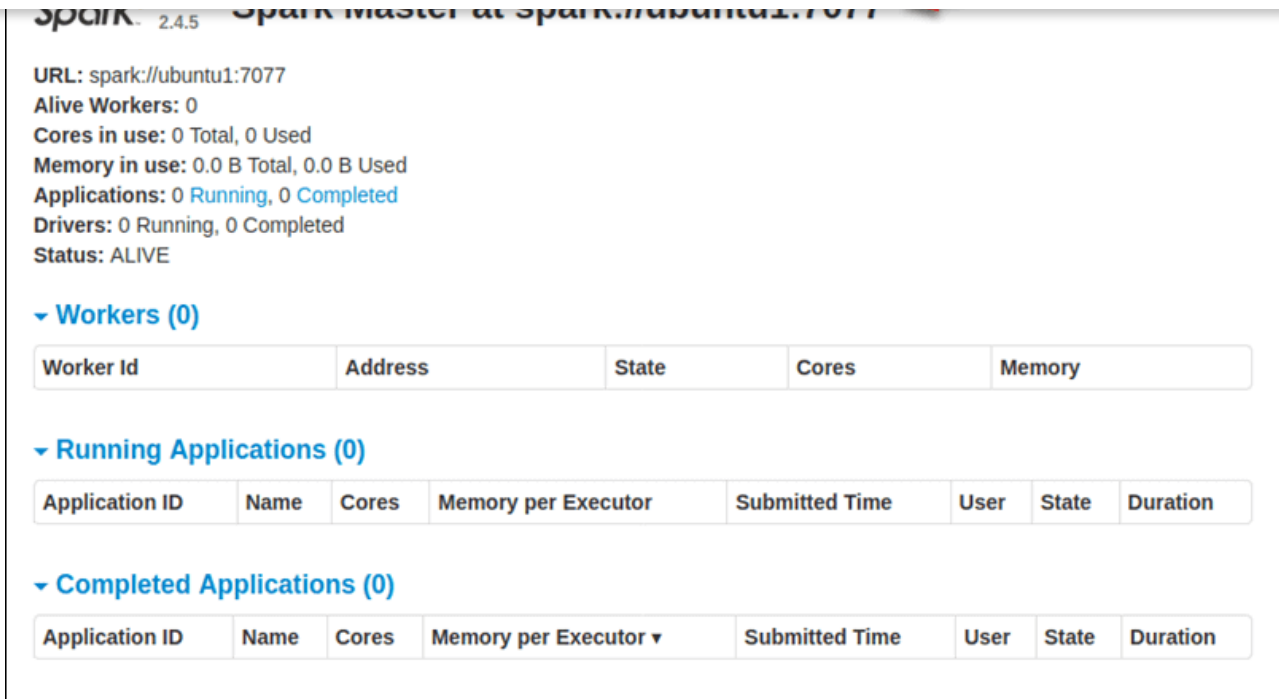
```
http://127.0.0.1:8080/
```



The page shows your **Spark URL**, status information for workers, hardware resource utilization, etc.

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE



The screenshot shows the Spark Master Web UI for a Spark 2.4.5 cluster. The URL is spark://ubuntu1:7077. The status is ALIVE. There are 0 alive workers, 0 cores in use, and 0.0 B memory in use. There are 0 running applications and 0 completed applications. Below the status, there are three expandable sections: Workers (0), Running Applications (0), and Completed Applications (0). Each section has a table with columns for Worker Id, Address, State, Cores, Memory, Application ID, Name, Cores, Memory per Executor, Submitted Time, User, State, and Duration.

The URL for Spark Master is the name of your device on port 8080. In our case, this is **ubuntu1:8080**. So, there are three possible ways to load Spark Master's Web UI:

1. 127.0.0.1:8080
2. localhost:8080
3. `deviceName:8080`



**Note:** Learn how to automate the deployment of Spark clusters on Ubuntu servers by reading our [Automated Deployment Of Spark Cluster On Bare Metal Cloud](#) article.

## Start Spark Slave Server (Start a Worker Process)

In this single-server, standalone setup, we will start one slave server along with the master server.

To do so, run the following command in this format:

```
start-slave.sh spark://master:port
```



## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

```
start-slave.sh spark://ubuntu1:7077
```

```
test@ubuntu1:~$ start-slave.sh spark://ubuntu1:7077
starting org.apache.spark.deploy.worker.Worker, logging to /opt/spark/logs/spark-test-or
g.apache.spark.deploy.worker.Worker-1-ubuntu1.out
test@ubuntu1:~$
```

Now that a worker is up and running, if you reload Spark Master's Web UI, you should see it on the list:

### Spark Master at spark://ubuntu1:7077

URL: spark://ubuntu1:7077

Alive Workers: 1

Cores in use: 2 Total, 0 Used

Memory in use: 1024.0 MB Total, 0.0 B Used

Applications: 0 Running, 0 Completed

Drivers: 0 Running, 0 Completed

Status: ALIVE

#### Workers (1)

Worker Id	Address	State	Cores	Memory
<a href="#">worker-20200331204050-10.0.2.15-46309</a>	10.0.2.15:46309	ALIVE	2 (0 Used)	1024.0 MB (0.0 B Used)

## Specify Resource Allocation for Workers

The default setting when starting a worker on a machine is to use all available CPU cores. You can specify the number of cores by passing the `-c` flag to the `start-slave` command.

For example, to start a worker and assign only **one CPU core** to it, enter this command:

```
start-slave.sh -c 1 spark://ubuntu1:7077
```

Reload Spark Master's Web UI to confirm the worker's configuration.

#### Workers (1)

Worker Id	Address	State	Cores	Memory
<a href="#">worker-20200401122203-10.0.2.15-33497</a>	10.0.2.15:33497	ALIVE	1 (0 Used)	1024.0 MB (0.0 B Used)

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

number. For gigabytes, use G and for megabytes, use M.

For example, to start a worker with 512MB of memory, enter this command:

```
start-slave.sh -m 512M spark://ubuntu1:7077
```

Reload the Spark Master Web UI to view the worker's status and confirm the configuration.

▼ Workers (1)				
Worker Id	Address	State	Cores	Memory
worker-20200401105553-10.0.2.15-43843	10.0.2.15:43843	ALIVE	2 (0 Used)	512.0 MB (0.0 B Used)

## Test Spark Shell

After you finish the configuration and start the master and slave server, test if the Spark shell works.

Load the shell by entering:

```
spark-shell
```

You should get a screen with notifications and Spark information. Scala is the default interface, so that shell loads when you run *spark-shell*.

The ending of the output looks like this for the version we are using at the time of writing this guide:



**VALENTINE'S DAY SPECIAL: Up to 33% OFF** selected dedicated servers and components!

GET COUPON CODE

```
/_/_/ ._/\/,._/ /_/\ \ version 2.4.5
/_/

Using Scala version 2.11.12 (OpenJDK 64-Bit Server VM, Java 11.0.6)
Type in expressions to have them evaluated.
Type :help for more information.

scala> █
```

Type **:q** and press **Enter** to exit Scala.

## Test Python in Spark

If you do not want to use the default Scala interface, you can switch to Python.

Make sure you quit Scala and then run this command:

```
pyspark
```

The resulting output looks similar to the previous one. Towards the bottom, you will see the version of Python.

```
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Welcome to

/_/_/ ._/\/,._/ /_/\ \ version 2.4.5
/_/

Using Python version 3.6.9 (default, Nov 7 2019 10:44:02)
SparkSession available as 'spark'.
>>>
```

To exit this shell, type **quit()** and hit **Enter**.

## Basic Commands to Start and Stop Master Server and Workers

Below are the basic commands for starting and stopping the Apache Spark master server and workers. Since this setup is only for one machine, the scripts you run default to the localhost.

## VALENTINE'S DAY SPECIAL: Up to 33% OFF selected dedicated servers and components!

GET COUPON CODE

```
start-master.sh
```

To **stop the master** instance started by executing the script above, run:

```
stop-master.sh
```

To **stop a running worker** process, enter this command:

```
stop-slave.sh
```

The Spark Master page, in this case, shows the worker status as DEAD.

Workers (2)				
Worker Id	Address	State	Cores	Memory
worker-20200331183244-10.0.2.15-45371	10.0.2.15:45371	DEAD	2 (0 Used)	1024.0 MB (0.0 B Used)
worker-20200331203427-10.0.2.15-37971	10.0.2.15:37971	ALIVE	2 (0 Used)	1024.0 MB (0.0 B Used)

You can **start both master and server** instances by using the start-all command:

```
start-all.sh
```

Similarly, you **can stop all instances** by using the following command:

```
stop-all.sh
```

## Conclusion

This tutorial showed you **how to install Spark on an Ubuntu machine**, as well as the necessary dependencies.

The setup in this guide enables you to perform basic tests before you start configuring a Spark cluster and performing advanced actions.